



HOW TO CRUNCH NUMBERS ON GPU

for fun and profit (with CUDA, C and R)

“It all started with video games”

2013 Video Game Market: \$93 Billion
2013 Business Analytics Market: \$14 Billion



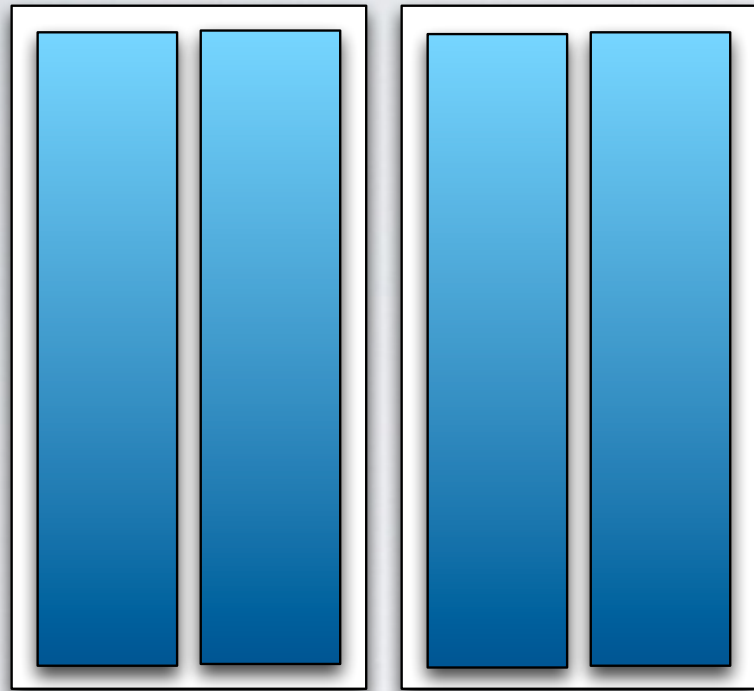
SHADERS





```
float4 Saturation : register(c0);
sampler2D implicitInputSampler : register(S0);

float4 main(float2 uv : TEXCOORD) : COLOR
{
    float3 LuminanceWeights = float3(0.299,0.587,0.114);
    float4 srcPixel = tex2D(implicitInputSampler, uv);
    float4 luminance = dot(srcPixel,LuminanceWeights);
    float4 dstPixel = lerp(luminance,srcPixel,Saturation);
    dstPixel.a = srcPixel.a;
    return dstPixel;
}
```

CPU vs GPU



Few cores




Hundreds cores



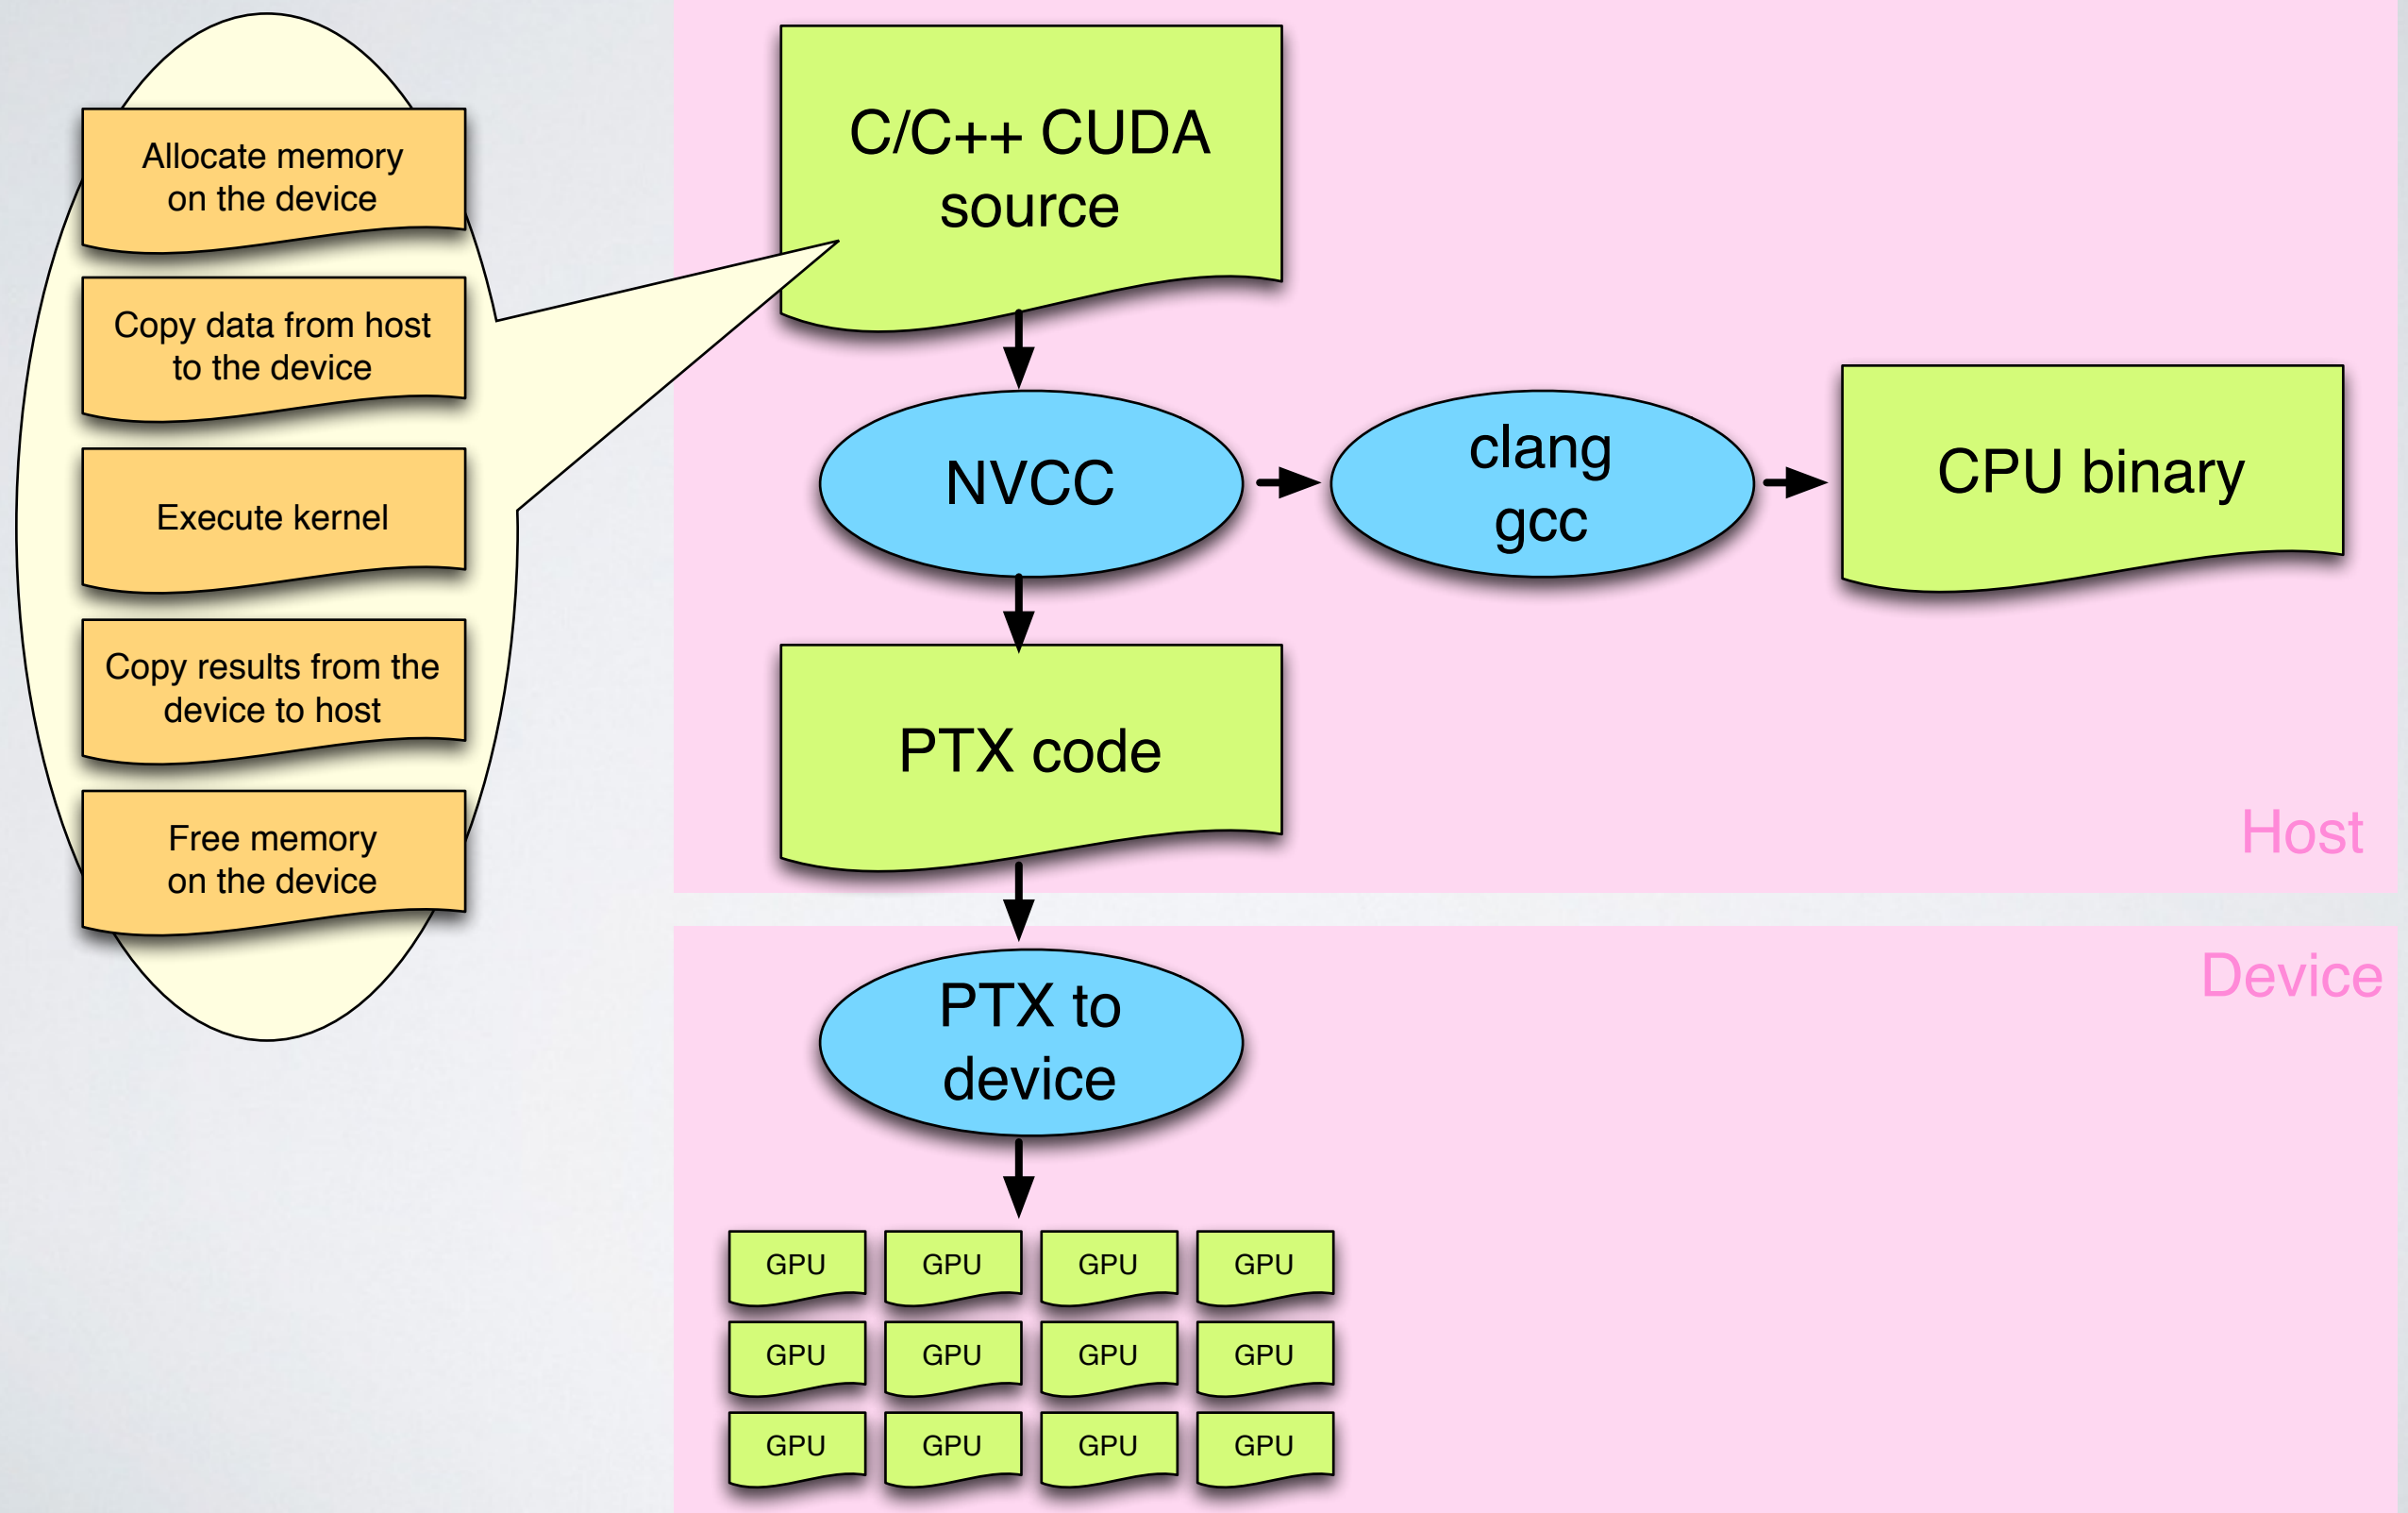
CUDA

“COMPUTE UNIFIED DEVICE ARCHITECTURE”



NVIDIA Corporation
ver 1.0: 2007
ver 6.0: Apr 2014

TYPICAL PROGRAM



SIMPLE SAMPLE

```
#include <cuda.h>
#define N 10

__global__ void kernel(int *a, int *b, int *c) {
    int tid = blockIdx.x;
    if (tid < N) {
        c[tid] = a[tid] + b[tid];
    }
}

int main() {
    int a[N], b[N], c[N];
    int *dev_a, *dev_b, *dev_c;

    cudaMalloc((void**)&dev_a, N * sizeof(int));
    cudaMalloc((void**)&dev_b, N * sizeof(int));
    cudaMalloc((void**)&dev_c, N * sizeof(int));

    for (int i = 0; i < N; ++i) {
        a[i] = -i;
        b[i] = i * i;
    }
```

```
    cudaMemcpy(dev_a, a, N * sizeof(int),
               cudaMemcpyHostToDevice);
    cudaMemcpy(dev_b, b, N * sizeof(int),
               cudaMemcpyHostToDevice);

    kernel<<<N, 1>>>(dev_a, dev_b, dev_c);

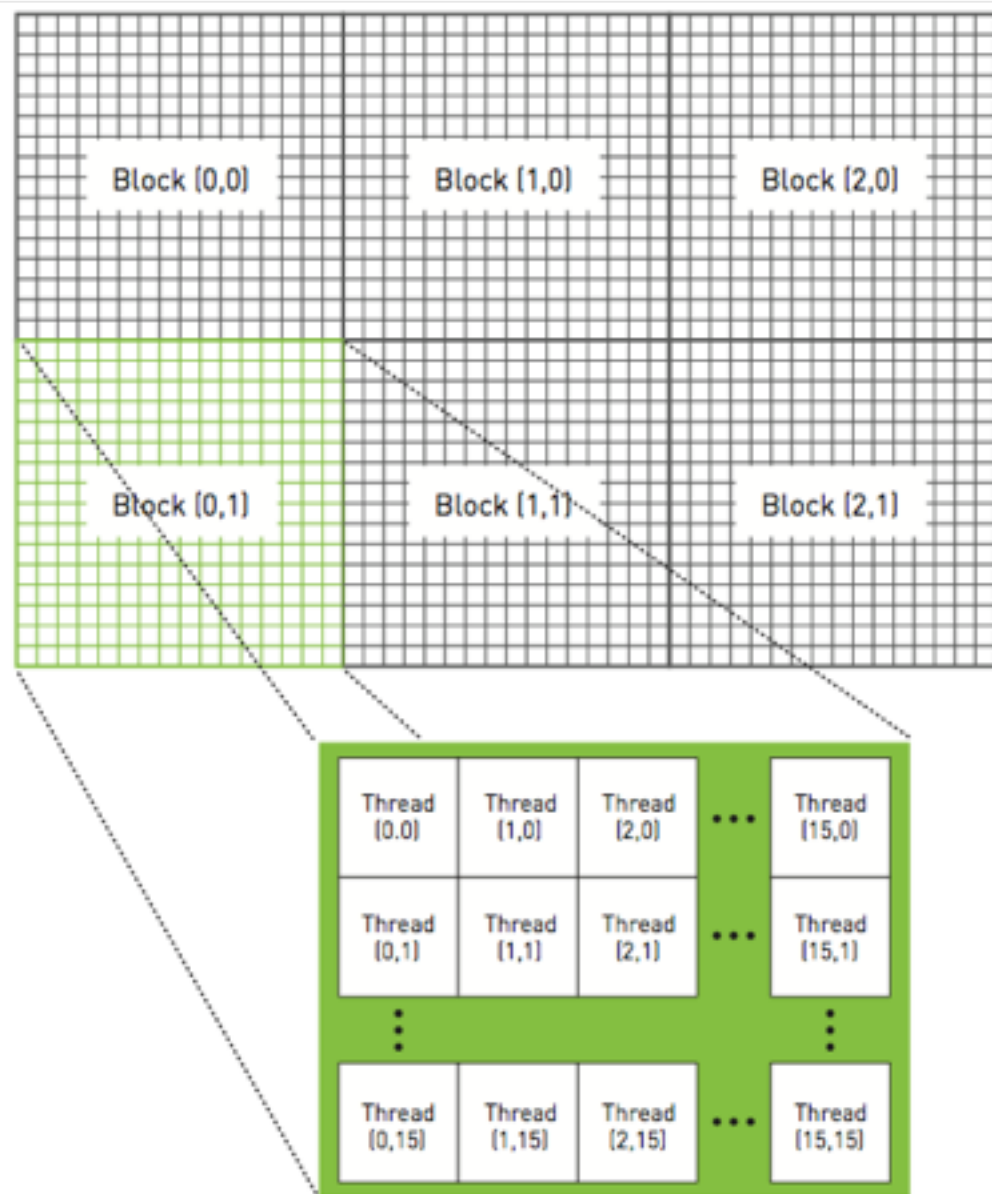
    cudaMemcpy(c, dev_c, N * sizeof(int),
               cudaMemcpyDeviceToHost);

    for (int i = 0; i < N; ++i) {
        printf("%d + %d = %d\n", a[i], b[i], c[i]);
    }

    cudaFree(dev_a);
    cudaFree(dev_b);
    cudaFree(dev_c);

    return 0;
}
```


BLOCKS AND THREADS






```
dim3 blocks(256, 256);  
dim3 threads(16, 16);  
kernel<<<block, threads>>>(data);
```


DEMO: MANDELBROT SET

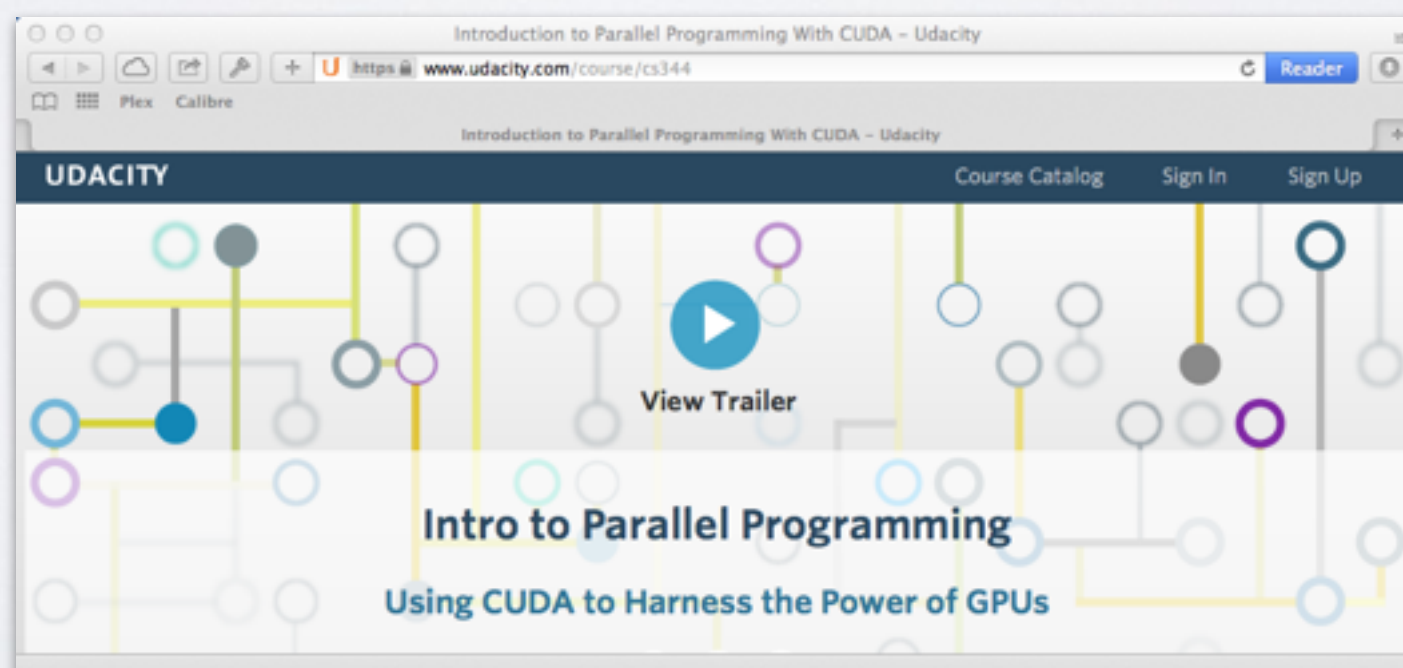




- Complex numbers
- Recurrent computation: $z_{n+1} = z_n^2 + c$
- c is the point of interest
- Start with $z_0 = 0+0i$
- c belongs to the set iff the absolute value of z_n remains bounded
- In practice: few hundred iterations
- If absolute value ≥ 2 then z_n will not be bounded

THINK PARALLEL

```
int sum_array(int[] a, int N) {  
    int sum = 0;  
    for (int i = 0; i < N; ++i) {  
        sum += a[i];  
    }  
    return sum;  
}
```



HUNGRY FOR MORE ?



<https://github.com/tvarkoj/project-mandel>